Response to the Reviewers — ICDEC-2026, Paper ID 6

We thank both reviewers. The camera-ready version uses the Springer LNNS template and addresses every comment below. We also tightened the evaluation (time-ordered hold-out, shortcut features removed, three seeds), so the numbers differ from the submitted version.

Reviewer 1

Organisation. The paper is restructured, and the dataset is now part of the results section (Sect. 4.1).

Table 1 not discussed. All tables and figures are now discussed in the text.

Architecture diagram. Added as Fig. 1 and explained in Sect. 1 and 3.

SHAP results and missing details. Fig. 2 and Sect. 3.1 show the SHAP ranking and how 44 features are kept. Table 1 lists all settings.

Figure 3 vs. conclusion. The old figure is replaced. All numbers now come from the same results, so the text, tables and conclusion agree.

Literature table. Table 3 compares related work, and ten of these methods are re-run on our data (Tables 4 and 8).

Reviewer 2

1. Multiclass vs. unseen classes. Clarified in Sect. 3.3: the classifier is multiclass on the training classes, and for unseen families any non-Benign prediction counts as a detected attack. Unseen families are evaluated by detection rate and false-positive rate, not accuracy.

2. "Zero-day". Replaced by "unseen attack families" throughout.

3. More evidence. Added class counts (Table 2), per-class precision, recall and F1, balanced accuracy, ROC-AUC, PR-AUC and a confusion matrix (Table 5), and results over three seeds (Table 4).

4. Feature-count ablation. Added as Table 6 for 10, 15, 23, 30, 40, 44 and all 64 usable features, with F1 and compute cost.

5. Inference time. The protocol is now stated (Sect. 4.2): 100,000 flows, batch size 1, one CPU thread, full pipeline from features to label. It takes 0.05 ms per flow.

Typos and grammar. The paper is rewritten and proofread.